

针对时间序列分类（TSC）的深度模型

Deep Learning for Time Series Classification¹

Xinyi Li

2019 年 11 月 28 日

¹Hassan Ismail Fawaz et al. “Deep learning for time series classification: a review”. In: *Data Mining and Knowledge Discovery* 33.4 (2019), pp. 917–963.

Overview

问题背景描述

Strong Baseline

深度模型结构

MLP/DNN

FCN

ResNet

Encoder

MCNN

Time Le-Net

MCDCNN

Time-CNN

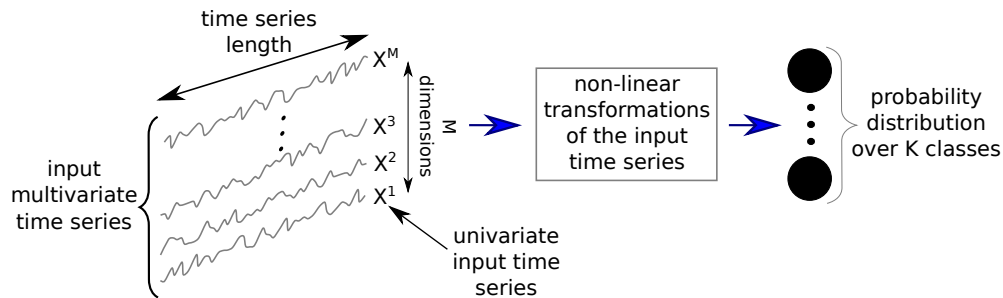
TWIESN

结论

Visualization Reasoning

TSC 社区生态

Problem Description



Different Learning Tasks

Univariate/Multi-variate/Panel

Univariate

1. Forecasting (e.g. fbprophet)
2. Regression
3. Classification

Multivariate

MTS

- ▶ different measurements of the **same** instance
- ▶ **high correlation**
- ▶ feeding features

panel data

- ▶ the same measurements on **different** instances
- ▶ i.i.d. assumption
- ▶ feeding sku/store/...

HIVE-COTE²:

SOTA classic algorithm³: Collective of Transformation based **Ensembles** (COTE) with a Hierarchical Vote system

1. Elastic Ensemble(**EE**): combination of 1-NN classifiers using different measurements
2. Shapelet Transform Ensemble(**ST**): top k shapelets (independent phase short pattern)
3. Bag-of-SFA-Symbols (**BOSS**) Ensemble: shapelets based on presence or absence
4. Time Series Forest (**TSF**): trained on selected $3\sqrt{m}$ features
5. Random Interval Features (**RIF**): spectral component of *Flat-COTE*

²Jason Lines, Sarah Taylor, and Anthony Bagnall. “Hive-cote: The hierarchical vote collective of transformation-based ensembles for time series classification”. In: *2016 IEEE 16th international conference on data mining (ICDM)*. IEEE. 2016, pp. 1041–1046.

³Anthony Bagnall et al. “The great time series classification bake off: a review and experimental evaluation of recent algorithmic advances”. In: *Data Mining and Knowledge Discovery* 31.3 (2017), pp. 606–660.

HIVE-COTE Algorithm⁴

ensemble of ensembles: ref to Flat-COTE

Algorithm 1 ProportionalEnsemble(*classifiers*, *train*, *test*)

```
1: trainAccs =  $\emptyset$ ;  
2: for  $i \leftarrow 1$  to  $|classifiers|$  do  
3:   trainAccsi = loocv(train, classifiers[i])  
4:   classifiersi.buildClassifier(train)  
5: testPreds =  $\emptyset$   
6: for  $i \leftarrow 1$  to  $|test|$  do  
7:   votes =  $\emptyset$   
8:   bsfWeight = -1;  
9:   bsfClass = -1;  
10:  for  $c \leftarrow 1$  to  $|classifiers|$  do  
11:     $p = classifiers_c.classify(test_i)$   
12:    votesp = votesp + trainAccsc;  
13:    if votesp > bsfWeight then  
14:      bsfWeight = votesp  
15:      bsfClass = p  
16:    testPredsi = bsfClass  
17: return testPreds
```

⁴<https://github.com/TonyBagnall/py-hive-cote>

MLP/DNN

Multi Layer Perceptrons/Fully-Connected(FC) Network

The Simplest DNN (e.g. `keras.Layers.Dense`)

$$\mathbf{X}_{i+1} = \sigma(\mathbf{W}_i \mathbf{X}_i + b_i)$$

Final (l -th) layer activate function: softmax

$$\hat{y}_k(\mathbf{X}_{l-1}) = (e^{\mathbf{W}_k \mathbf{X}_{l-1} + b_k}) / (\sum_{i=1}^K e^{\mathbf{W}_i \mathbf{X}_{l-1} + b_i})$$

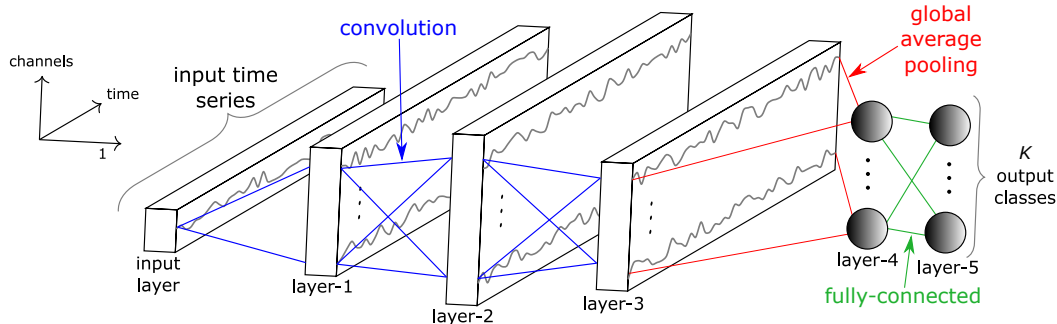
Objective loss: categorical cross entropy

$$Loss(\mathbf{X}) = - \sum_{i=1}^K y_i \log \hat{y}_i$$

minimized to learn the weights using **gradient descent** method

FCN⁵

Fully Convolutional Neural Network




convolution layer (\forall time stamp t shares filter ω with length l)

$$\mathbf{C}_t = \sigma(\omega * \mathbf{X}_{t-l/2:t+l/2} + b) | \forall t \in [1, T]$$

⁵ John Cristian Borges Gamboa. "Deep learning for time-series analysis". In: *arXiv preprint arXiv:1701.01887* (2017).

ResNet⁶

Residual Network




⁶Zhiguang Wang, Weizhong Yan, and Tim Oates. "Time series classification from scratch with deep neural networks: A strong baseline". In: *2017 international joint conference on neural networks (IJCNN)*. IEEE. 2017, pp. 1578–1585.

Encoder⁷

Hybrid deep CNN based on FCN

Modified from FCN

- ▶ GAP layer → attention layer (**careful design for pre-train**)
- ▶ normalization for each Conv layer output:
 1. ReLU → PReLU activation function (+ parameter)
 2. + dropout regularization
 3. + max pooling



⁷Joan Serrà, Santiago Pascual, and Alexandros Karatzoglou. “Towards a Universal Neural Network Encoder for Time Series.”. In: *CCIA*. 2018, pp. 120–129.

MCNN⁸

Multi-scale Convolutional Neural Network

Similar to Traditional CNN:

1. 2 Conv layer (with max pooling)
2. 1 Fully-connected layer
3. final softmax layer

Heavy data pre-preprocessing step:

Window Slicing(WS)

for data augmentation:

1. slide a window over raw input
2. extract subsequences

Before training, \forall subsequence


Transformations (parallel)

1. identity mapping
 - ▶ keep unchanged \rightarrow 1-st Conv
2. down-sampling
 - ▶ \rightarrow different shorter lengths subsequences
 - ▶ \rightarrow 1-st Conv
3. smoothing:
 - ▶ equal length one
 - ▶ \rightarrow 1-st Conv
 - ▶ \rightarrow 2-nd Conv

⁸Zhicheng Cui, Wenlin Chen, and Yixin Chen. "Multi-scale convolutional neural networks for time series classification". In: *arXiv preprint arXiv:1603.06995* (2016).

MCNN

Framework



Time Le-Net⁹

Inspired by Le-Net¹⁰, like CNN: 2 Conv + FC + final softmax

Compared to FCN:

GAP \rightarrow FC

Local max pooling

- ▶ take max in a local pooling
- ▶ + invariance to small perturbations
- ▶ shorten a time series

Still #parameters \uparrow #invariance \downarrow

Data augmentation to prevent **overfitting**
especially on relatively small datasets

Window Slicing (WS)

= method used in MCNN

Window Warping (WW)

For a time series with length l

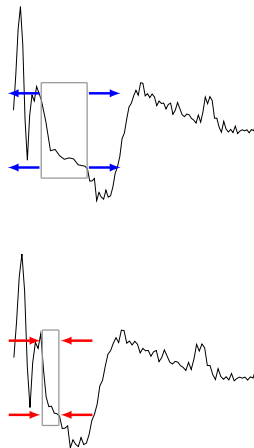
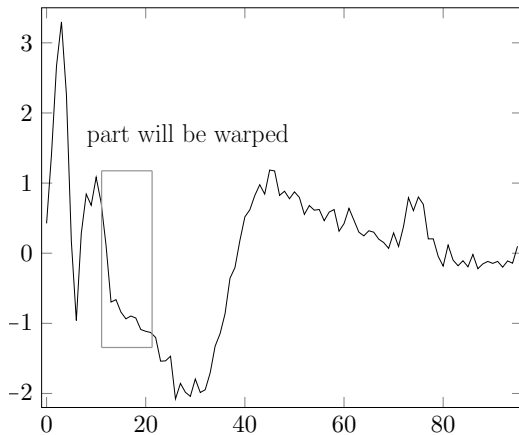
1. dilate ($\times 2$) $\rightarrow 2l$
2. squeeze ($\times \frac{1}{2}$) $\rightarrow \frac{1}{2}l$

⁹Arthur Le Guennec, Simon Malinowski, and Romain Tavenard. “Data augmentation for time series classification using convolutional neural networks”. In: 2016.

¹⁰Yann LeCun et al. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324.

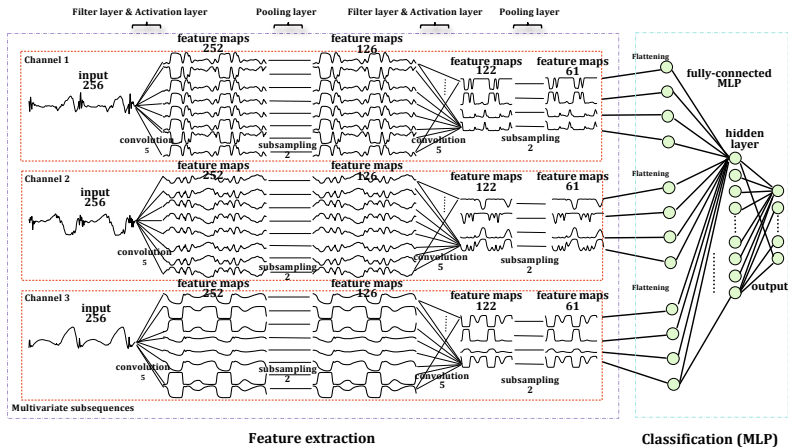
Window Warping (WW)

Then feed all data (length= $l, 2l, \frac{1}{2}l$) into network



MCDCNN¹¹

Multi Channel Deep Convolutional Neural Network: **Independent** Conv specified for **MTS**



¹¹Yi Zheng et al. "Time series classification using multi-channels deep convolutional neural networks". In: *International Conference on Web-Age Information Management*. Springer. 2014, pp. 298–310.

Time-CNN¹²

both for univariate and multivariate

Main differences compared to previous models:

1. loss function: categorical cross-entropy \rightarrow MSE
2. activate function of final layer: softmax \rightarrow sigmoid

$$\sum_{i=1}^K p(\hat{Y}_i) \neq 1$$

3. throughout CNN: local max pooling \rightarrow local average pooling

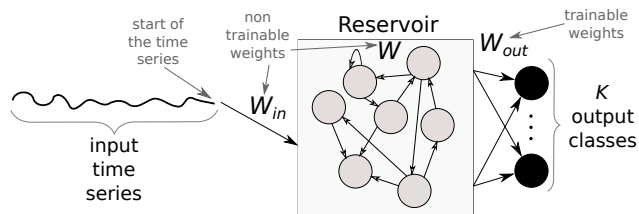
Compared to some models:

- ▶ apply 1 Conv for all dimensions (unlike MCD CNN)
- ▶ Conv directly fully connected to final layer (modify FCN: GAP \rightarrow FC)

¹²Bendong Zhao et al. “Convolutional neural networks for time series classification”. In: *Journal of Systems Engineering and Electronics* 28.1 (2017), pp. 162–169.

Echo State Network (ESN)¹³

Based on RNN



Core part: reservoir

- sparsely connected random RNN
- each neuron create its own nonlinear activation of the incoming signal

Internal state at time t : $I(t) \in R^{N_r}$, $N_r = \# \text{neurons inside reservoir}$

$$I(t) = \sigma(W_{in}X(t) + WI(t-1)) | \forall t \in [1, T]$$

Finally

$$\hat{Y}(t) = W_{out}I(t)$$

¹³Claudio Gallicchio and Alessio Micheli. "Deep echo state network (deepesn): A brief survey". In: *arXiv preprint arXiv:1712.04323* (2017).

TWIESN¹⁴

Time Warping Invariant Echo State Network

ESN originally proposed for time series forecasting → predicts a probability distribution

Training

1. reservoir space ($\forall t$): $X(t) \rightarrow$ **higher** dimensional space
2. train a Ridge classifier to predict the class of each $X(t)$

Testing

1. $\forall X(t)$, trained Ridge classifier predicts a probability distribution
2. for each class k , calculate a posteriori probability
3. label k with the largest average probability

$$\arg \max_k \frac{1}{T} \sum_{t=1}^T \hat{Y}_k(t)$$

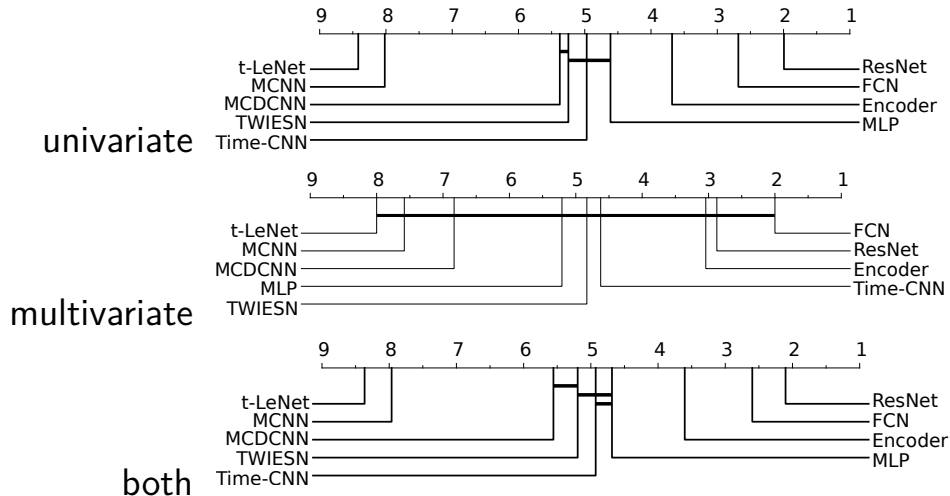
¹⁴Pattreeya Tanisaro and Gunther Heidemann. "Time series classification using time warping invariant echo state networks". In: *2016 15th IEEE International Conference on Machine Learning and Applications (ICMLA)*. IEEE. 2016, pp. 831–836.

Summary

Methods	Architecture							
	#Layers	#Conv	#Invar	Normalize	Pooling	Feature	Activate	Regularize
MLP	4	0	0	None	None	FC	ReLU	Dropout
FCN	5	3	4	Batch	None	GAP	ReLU	None
ResNet	11	9	10	Batch	None	GAP	ReLU	None
Encoder	5	3	4	Instance	Max	Att	PReLU	Dropout
MCNN	4	2	2	None	Max	FC	Sigmoid	None
t-LeNet	4	2	2	None	Max	FC	ReLU	None
MCDCNN	4	2	2	None	Max	FC	ReLU	None
Time-CNN	3	2	2	None	Avg	Conv	Sigmoid	None

Overall Performance

Critical Difference Diagram (dataset: UCR/UEA)



For Different Datasets

Themes (#)	MLP	FCN	ResNet	Encoder	MCNN	t-LeNet	MCDCNN	Time-CNN	TWIESN
DEVICE (6)	0.0	50.0	83.3	0.0	0.0	0.0	0.0	0.0	0.0
ECG (7)	14.3	71.4	28.6	42.9	0.0	0.0	14.3	0.0	0.0
IMAGE (29)	6.9	34.5	48.3	10.3	0.0	0.0	6.9	10.3	0.0
MOTION (14)	14.3	28.6	71.4	21.4	0.0	0.0	0.0	0.0	0.0
SENSOR (16)	6.2	37.5	75.0	31.2	6.2	6.2	6.2	0.0	12.5
SIMULATED (6)	0.0	33.3	100.0	33.3	0.0	0.0	0.0	0.0	0.0
SPECTRO (7)	14.3	14.3	71.4	0.0	0.0	0.0	0.0	28.6	28.6

Length	MLP	FCN	ResNet	Encoder	MCNN	t-LeNet	MCDCNN	Time-CNN	TWIESN
<81	5.43	3.36	2.43	2.79	8.21	8.0	3.07	3.64	5.5
81-250	4.16	1.63	1.79	3.42	7.89	8.32	5.26	4.47	5.53
251-450	3.91	2.73	1.64	3.32	8.05	8.36	6.0	4.68	4.91
451-700	4.85	2.69	1.92	3.85	7.08	7.08	5.62	4.92	4.31
701-1000	4.6	1.9	1.6	3.8	7.4	8.5	5.2	6.0	4.5
>1000	3.29	2.71	1.43	3.43	7.29	8.43	4.86	5.71	6.0

Train size	MLP	FCN	ResNet	Encoder	MCNN	t-LeNet	MCDCNN	Time-CNN	TWIESN
<100	4.3	2.03	1.67	4.13	7.67	7.73	6.1	4.37	4.77
100-399	4.85	2.76	2.06	3.24	7.71	8.12	4.59	4.97	4.5
400-799	3.62	2.38	1.75	3.5	8.0	8.62	4.38	5.0	5.88
>799	3.85	2.85	1.62	2.08	7.92	8.69	4.62	4.85	6.92

Class Activation Map (CAM)¹⁵

GAP layer + softmax

univariate time series $A_m(t)$ produced by m -th filter in Last Conv, input to neuron of class c

$$\begin{aligned} z_c &= \sum_m w_m^c \sum_t A_m(t) \\ &= \sum_m \sum_t w_m^c A_m(t) \end{aligned}$$

Therefore, CAM_c that explains the classification as label c

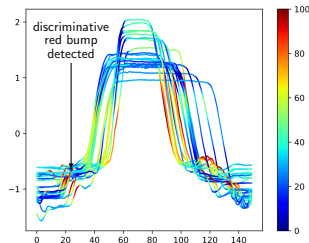
$$CAM_c(t) = \sum_m w_m^c A_m(t)$$

constructed a new time series for further observation and explanation.

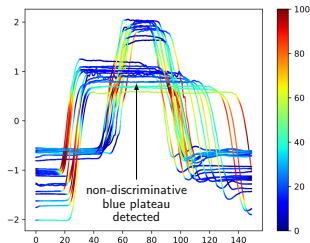
¹⁵Bolei Zhou et al. "Learning deep features for discriminative localization". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, pp. 2921–2929.

Find discriminative Region

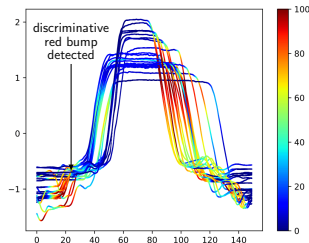
Example: FCN/ResNet on GunPoint dataset (achieve 100% accuracy)



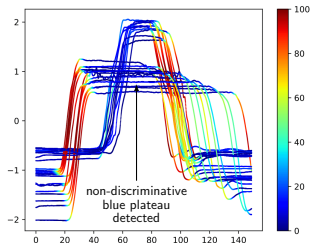
(a) FCN on GunPoint: Class-1



(b) FCN on GunPoint: Class-2



(c) ResNet on GunPoint: Class-1



(d) ResNet on GunPoint: Class-2

Libraries/Implements/Community

sktime¹⁶ & its extensions

Sktime

- ▶ based on classic models (shallow)
- ▶ scikit-learn interface compatible

Sktime-dl

- ▶ use Keras to implement all 9 deep models above
- ▶ 暂时不能直接安装 (MacOS)


UEA & UCR Time Series Classification Repository

- ▶ 128 TSC datasets + 30 MTS datasets
- ▶ Collect a bunch of **classic** algorithms

¹⁶Markus Löning et al. *sktime: A Unified Interface for Machine Learning with Time Series*. 2019. eprint: arXiv:1909.07872.

Thanks

All codes, slides and papers available

 [li-xin-yi/deep_time_series_share_slide](https://github.com/li-xin-yi/deep_time_series_share_slide)